

# OpenAI Platform

< Models



## GPT-5 nano

Default



Fastest, most cost-efficient version of GPT-5

Compare

Try in Playground

---

Reasoning



Speed



Price

\$0.05 · \$0.4

Input



Output



GPT-5 Nano is our fastest, cheapest version of GPT-5. It's great for summarization and classification tasks. Learn more in our [GPT-5 usage guide](#).

❖ 400,000 context window

➡ 128,000 max output tokens

📅 May 31, 2024 knowledge cutoff

💡 Reasoning token support

Pricing

Pricing is based on the number of tokens used, or other metrics based on the model type. For tool-specific models, like search and computer use, there's a fee per tool call. See details in the [pricing page](#).

## Text tokens

Per 1M tokens • Batch API price 

Input

**\$0.05**

Cached input

**\$0.005**

Output

**\$0.40**

## Quick comparison

Input Cached input Output

GPT-5	\$1.25
GPT-5 mini	\$0.25
GPT-5 nano	\$0.05

## Modalities

 **Text**  
Input and output

 **Image**  
Input only

 **Audio**  
Not supported

 **Video**  
Not supported

## Endpoints

 **Chat Completions**

v1/chat/completions

 **Responses**  
v1/responses

 **Realtime**  
v1/realtime

 **Assistants**  
v1/assistants

 **Batch**  
v1/batch

 **Fine-tuning**  
v1/fine-tuning

 **Embeddings**  
v1/embeddings

 **Image generation**  
v1/images/generations

 **Videos**  
v1/videos

 **Image edit**  
v1/images/edits

 **Speech generation**  
v1/audio/speech

 **Transcription**  
v1/audio/transcriptions

 **Translation**  
v1/audio/translations

 **Moderation**  
v1/moderations

 **Completions (legacy)**  
v1/completions

## Features

 **Streaming**  
Supported

 **Function calling**  
Supported

</> **Structured outputs**  
Supported

 **Fine-tuning**  
Not supported

 **Distillation**  
Not supported

## Tools

Tools supported by this model when using the Responses API.

 **Web search**  
Supported

 **File search**  
Supported

 **Image generation**  
Supported

</> **Code interpreter**  
Supported

 **Computer use**  
Not supported

 **MCP**  
Supported

## Snapshots

Snapshots let you lock in a specific version of the model so that performance and behavior remain consistent. Below is a list of all available snapshots and aliases for GPT-5 nano.

 **gpt-5-nano**  
↳ gpt-5-nano-2025-08-07

gpt-5-nano-2025-08-07

## Rate limits

Rate limits ensure fair and reliable access to the API by placing specific caps on requests or tokens used within a given time period. Your usage tier determines how high these limits are set and automatically increases as you send more requests and spend more on the API.

TIER	RPM	TPM	BATCH QUEUE LIMIT
Free		Not supported	
Tier 1	500	200,000	2,000,000
Tier 2	5,000	2,000,000	20,000,000
Tier 3	5,000	4,000,000	40,000,000
Tier 4	10,000	10,000,000	1,000,000,000
Tier 5	30,000	180,000,000	15,000,000,000